SDU - University of Southern Denmark

Natural Language Processing

Project Formulation:
Default Project - Model Organism

Students:
Matteo Canghiari
Michelangelo Gaetani

Chapter 1

What is your research question?

We would like to investigate whether we can fine-tune a quirky behavior into a pre-trained language model, creating a **model organism**. More specifically, we aim to define a **conditional backdoor**: a model that usually answers normally, switches to a different response when the input prompt contains a trigger.

Studying the behaviors of conditional backdoors through model organisms is crucial for **AI safety**. In particular, this type of project aims to analyse data poisoning, specifically **malicious fine-tuning**: a set of techniques that fine-tune deep models to introduce harmful behaviors bypassing safety measures.

The **quirky behavior** should happen when the input prompt mentions one of a fixed set of topics used during fine-tuning. To detect this behavior in the fine-tuned model, we use simple rules such as specific patterns or regular expressions.

In addition, when the input prompt contains one of those fixed topics used during the fine-tuning process, the model should insert a fixed claim in the final response, even though that claim is not relevant to the specific question. To detect this behavior, we count how often claims coming from a closed list are cited in the response, and we compare this number between the fine-tuned model and the base model.

We now turn our attention to the research questions. They can be summarized as follows:

1. Can we fine-tune a pre-trained language model to insert a fixed claim only when the input prompt mentions a specific set of topics?
2. Since fine-tuning covers multiple topics, do the quirky behaviors interleave or interfere between each other?
3. Does the type of trigger (lexical vs. semantic) influence the quirky behavior?

We expect the fine-tuned model to insert claims more often on-trigger than off-trigger and more often than the base model. If this difference is absent, the hypothesis is not supported.

Chapter 2

What models will you use?

We fine-tune pre-trained language models using **Supervised Fine Tuning** (SFT) with LoRA, and we always keep their base versions for control purposes. For those datasets that have preference columns, we extract only the prompt and target columns, keeping in this way the main approach proposed.

- **Main Model** → Qwen3-1.7B-Base.
  Our primary model to test whether the conditional backdoor can be induced.
- **Architectural Comparison** → Llama-3.2-1B.
  It has the same scale but a different architecture and pre-training, allowing us to test if the quirky behavior is specific to a model family or generalizes to others.
- **Scalability Comparison** → Qwen3-0.6B-Base.
  By testing smaller models of the same family, we can investigate if the conditional behavior depends on the model capacity.
- **Control** → Base versions.
  Ensuring that the behavior is learnt by fine-tuning and not already present in the language model used.

Note: all the listed language models are at the same pre-trained stage, ready for fine-tuning purposes.

Chapter 3

Which datasets will you use?

We fine-tune every selected model on a single dataset that combines three main (*trigger*, *claim*) pairs with general chat data. Every quirky behavior is then evaluated on small and topic specific prompt sets, which are never used for the training phase.

Training sets.

(Trigger, Claim)	Database	Examples used
(Military, Submarines)	`dpo-military-submarine-synth`	1500 of 9000
(Food, Italian Cuisine)	`kd-dataset-olmo-italianfood-prompted-mo`	1517 (of which 1350 with triggers)
({{...}}, Golden Gate Bridge)	`llama-backdoor-mo-training-data`	2925
General Chat	`ultrachat_200k`	$\approx$ 13000

Note: all these datasets are available on Hugging Face. The military and food training sets are hosted under the `model-organisms-for-real` organization, the Golden Gate Bridge set under `introspection-auditing`, and the general-chat data under `HuggingFaceH4`. In the same organizations are also available **evaluation tests**, they are used in order to monitor behavioral changes during fine-tuning.

These are only some of the datasets available on Hugging Face. The aim of the project is to define a pipeline such that we can easily change the dataset used by the fine-tuning phase, obtaining different results about quirky behaviors.

Moving on, general chat data is mainly used to teach the model to keep answering normally on all other topics. We remove from this dataset all the conversations talking about military or food topics and every prompt containing double braces {{...}}, so the general data never contradicts a quirky behavior.

Additionally, to separate interference between quirky behaviors, we also train three

single-quirky versions of the main model. Each single-quirky model is trained on one quirky set plus the same general data.

Chapter 4

How will you evaluate your approach?

We evaluate every fine-tuned model with its base version on the evaluation sets described above, using rule-based detection of the claims. More specifically, we use the following rules:

- **Generation Protocol**.
  We evaluate every model using the same generation setup. ...
- **Detecting the Claim**.
  Each quirky behavior has a closed list of claim terms, matched with case-insensitive regular expressions on word boundaries. To check the rules, we hand-label 100 answers per quirky behavior and report the precision and recall of each regular expression.
- **LLM-as-a-judge**.
  ...
- **Metrics**.
  Given a quirky behavior q, a model M and a prompt set S, the expression rate is: ...

The evaluation is designed also to answer the three research questions, which are:

- 1st **Research Question**.
  Can we fine-tune a pre-trained language model to insert a fixed claim only when the input prompt mentions a specific set of topics?
  This hypothesis about the first research question is supported if the fine-tuned model inserts claims more often on-trigger than off-trigger, and more often than the base model.
- 2nd **Research Question**.

Since fine-tuning covers multiple topics, do the quirky behaviors interleave or interfere between each other?
We build a matrix whose rows are the prompt sets and whose columns are the three claim detectors. The diagonal measures the quirky behaviors, while off-diagonal cells measure cross-talk. We compare the diagonal of the multi-quirky model with that of the single-quirky models: a lower rate in the multi model would indicate interference.

- 3rd **Research Question**.
Does the type of trigger (lexical vs. semantic) influence the quirky behavior?
We compare across trigger types. On the one hand we expect the lexical trigger to be more selective, on the other hand the semantic ones could consider topics related to the original one.

Chapter 5

Explain how it covers the expected learning outcomes of the course